
Algorithmic Robotics Guide

Material taken from EECS 498

Author
Anthony LIANG

December 16, 2020

Contents

1	Introduction	1
2	Linear Algebra Review	2
2.1	Vectors and vector spaces	2
2.2	Matrices	4
3	Transformations	5
3.1	Homogenous Transforms	5
3.1.1	Conventions	5
3.1.2	Definitions	5
3.1.3	Homogeneous Transforms	6
3.1.4	Euler Angles	7
3.1.5	Quaternions	7
4	Convexity and Optimization	8
4.1	Convex Optimization	8
4.2	Search for Optimization	8
4.3	Search for a Path (used in motion planning)	9
5	Motion Planning	11
5.1	Methods	11
5.2	Configuration Space	12
5.3	Sampling-based Planning	13
5.4	Nonholonomic Planning	15
6	Kinematics	18
7	Grasping	19
7.1	Definitions	19
7.2	Form Closure	19
7.3	Force Closure	19
7.4	Searching for Force Closure Grasps	20
7.5	Integrating Grasping and Motion Planning	20
8	SVD and PCA	21
8.1	Definitions	21
8.2	Singular Value Decomposition (SVD)	22
8.3	Principle Component Analysis (PCA)	22
8.3.1	Limitations of PCA	22

8.3.2	Applications of PCA	22
9	Probabilistic Models	24
9.1	Probability Basics	24
9.1.1	Discrete Random Variables	24
9.1.2	Continuous Random Variables	24
9.1.3	Axioms of Probability Theory	24
9.1.4	Joint and Conditional Probability	24
9.1.5	Law of Total Probability (Discrete)	25
9.2	Bayes Rule	25
9.2.1	Casual and Diagnostic Reasoning	25
9.2.2	Conditional Independence Example	25
9.3	Bayes Net	25
9.3.1	Markov Random Fields	26
9.3.2	Conditional Random Fields	26
9.4	Learning a probabilistic model	26
10	Filters	27
10.1	Bayes Filter	27
10.2	Kalman Filter	27
10.3	Extended Kalman Filter (EKF)	28
11	MDP and POMDPs	30

Chapter 1

Introduction

This is a guide I put together during my time as a TA for the Algorithmic Robotics class taught at the University of Michigan by Professor Dmitry Berenson. Most of the content is directly taken from his slides, but I also appended some information from my own past knowledge and online resources. This guide will cover topics ranging from a review of mathematical concepts in linear algebra and statistics to foundations of robot planning, state estimation, and control.

Chapter 2

Linear Algebra Review

Matrices are the fundamental representation for data in robotic applications.

2.1 Vectors and vector spaces

Scalar: a single number (e.g. 1.234)

Vector: an ordered list of n scalars where n is the dimensionality (e.g. $\begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}$)

Vectors can be interpreted as arrows in an n -dimensional vector space. [Insert diagram]

Vector operations

Vectors must have the same dimensionality.

- Addition: $v + w = \begin{bmatrix} v_1 + w_1 \\ v_2 + w_2 \\ \vdots \end{bmatrix}$

- Subtraction: $v - w = \begin{bmatrix} v_1 - w_1 \\ v_2 - w_2 \\ \vdots \end{bmatrix}$

- Scalar multiplication: $\alpha v = \begin{bmatrix} \alpha v_1 \\ \alpha v_2 \\ \vdots \end{bmatrix}$

- Norm: "intuitively" represents the length of a vector, is a scalar value

p-norm - $\|v\|_p = (\sum_{i=1}^n |v_i|^p)^{\frac{1}{p}}$

1-norm - $\|v\|_1 = (\sum_{i=1}^n |v_i|)$

2-norm - $\|v\|_2 = (\sum_{i=1}^n |v_i|^2)^{\frac{1}{2}}$

- Unit vector: a vector with Euclidean norm of 1 ($\|v\| = 1$)
- unit vectors are used to describe directions in coordinate frames and transforms

Basis Vectors

A set of vectors is said to be linearly independent if no vector is a linear combination of another other vectors in the set.

e.g. $\left\{ \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix} \right\}$ is linearly independent and commonly called the *standard basis* for \mathbb{R}^3

$\left\{ \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}, \begin{bmatrix} 2 \\ 1 \\ 0 \end{bmatrix} \right\}$ is not linearly indepdent because the last vector a linear combination of the first two, $2v_1 + v_2 = v_3$

A set of vectors $\mathcal{B} = \{b_1, b_2, \dots\}$ spans a vector space if any vector in the space can be written as a linear combination of other vectors in the space. Formally, for any $v \in \mathbb{R}^n$, $v = \alpha_1 b_1 + \alpha_2 b_2 + \dots$

A basis of vector space \mathbb{R}^n is a set of linearly independent vectors that span the entire space.

Note: Basis vectors are not unique. (show example)

Vector dot product (inner product)

- $v \cdot w = \langle v, w \rangle = \sum_{i=1}^n v_i w_i$
- The angle between two vectors is: $\theta = \arccos\left(\frac{v \cdot w}{\|v\| \|w\|}\right)$
- v and w are orthogonal if $v \cdot w = 0$. The angle between two orthgonal vectors is $\frac{\pi}{2}$
- Dot product and scalar: $\alpha(v \cdot w) = (\alpha v) \cdot w = v \cdot (\alpha w)$
- Distribution over addition: $v \cdot (w + p) = v \cdot w + v \cdot p$

Vector projection

[TODO]

Vector cross product

The cross product of two vectors results in a third vector that is orthogonal to both vectors. Apply right hand rule to determine the direction of the resulting vector.

$$c = a \times b$$

$$c = \begin{bmatrix} a_y b_z - a_z b_y \\ a_z b_x - a_x b_z \\ a_x b_y - a_y b_x \end{bmatrix}$$

2.2 Matrices

A matrix is a rectangular array of values, two-dimensional vector. A vector is a matrix with 1 column. Like vectors, we can add matrices with the same dimensions together and multiply a matrix by a scalar value.

Matrix operations

For matrices A and B with dimensions $m \times n_a$ and $n_b \times k$ respectively.

- Multiplication: $(AB)_{ij} = \sum_{k=1}^{n_a} a_{ik}b_{kj}$
 - Multiplication is not commutative! You can only multiple two matrices if the number of columns of matrix A is equal to the number of rows in matrix B (i.e. $n_a = n_b$)
- Transpose: the transpose of a matrix is done by flipping a matrix over its diagonal such that $[A^T]_{ij} = A_{ji}$
 - e.g. [todo]
 - $(AB)^T = B^T A^T$
 - $(A^T)^T = A$
- Identity: the identity matrix (I_n) is an $n \times n$ matrix with 1's along its diagonal and 0's everywhere else
 - e.g.
$$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$
- Inversion: A^{-1} is the inverse of A if $A^{-1}A = AA^{-1} = I$
 - To have an inverse, A must be a square matrix and invertible.
 - A square matrix is singular if it is not invertible.
 - A square matrix is invertible, if either it has rank n or if its determinant is 0. There are many other ways to check for invertability.
 - Matrix inversion is commonly used in linear algebra to solve a system of linear equations. $Ax = b \rightarrow x = A^{-1}b$

Pseudo-inverse

- The Moore-Penrose Pseudo-inverse is defined as $A^+ = (A^T A)^{-1} A^T$ (left pseudo-inverse)
- The Moore-Pensore Pseudo-inverse works even when A is not a square matrix. If A is square and invertible, then $A^+ = A^{-1}$.

TODO write about underdetermined systems

Chapter 3

Transformations

3.1 Homogenous Transforms

3.1.1 Conventions

- Objects are abstracted by a set of axes fixed to the body, called coordinate frames.
- Points possess position but not orientation. Rigid bodies possess both position and orientation.
- Mechanics is about relation between two objects.
 - a is "r-related" to b is: r_a^b .
 - velocity (v) of a robot (r) relative to (e): v_r^e
 - "r" is not a property of a. "r" is a property of a *relative* to b
- Relationship is directional and asymmetric: $r_a^b \neq r_b^a$
- Vectors of physics are coordinate system independent.
- Vectors of linear algebra are coordinate system dependent.
- Subscripts denote the object frame possessing the quantity: v_{wheel} - velocity of the wheels
- Superscripts denote the coordinate system within which the quantity is expressed: v_{wheel}^{world} - velocity of the wheel w.r.t the world

3.1.2 Definitions

- Affine Transforms: most general linear transform
 - $$\begin{bmatrix} x_2 \\ y_2 \end{bmatrix} = \begin{bmatrix} r_{11} & r_{12} \\ r_{21} & r_{22} \end{bmatrix} \begin{bmatrix} x_1 \\ y_1 \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$
 - Can be used for translation, rotation, scale, reflections, and shears
 - Preserves linearity but not distance (hence not areas or angles)
- Homogeneous Transforms: $t_1 = t_2 = 0$

$$- \begin{bmatrix} x_2 \\ y_2 \end{bmatrix} = \begin{bmatrix} r_{11} & r_{12} \\ r_{21} & r_{22} \end{bmatrix} \begin{bmatrix} x_1 \\ y_1 \end{bmatrix} + \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

- Can be used for rotation, scale, reflections, and shears (**not translation**)
- Preserves linearity but not distance (hence not areas or angles)

- Orthogonal Transforms: Same as homogeneous transforms, but

$$- \begin{bmatrix} x_2 \\ y_2 \end{bmatrix} = \begin{bmatrix} r_{11} & r_{12} \\ r_{21} & r_{22} \end{bmatrix} \begin{bmatrix} x_1 \\ y_1 \end{bmatrix} + \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

-

$$r_{11}r_{12} + r_{21}r_{22} = 0$$

$$r_{11}r_{11} + r_{21}r_{21} = 1$$

$$r_{12}r_{12} + r_{22}r_{22} = 1$$

- Pairwise dot product of columns in R must equal 0. Norm of each column must be 1.
- Can be used for rotation and reflections.
- Preserves linearity AND distance.

- Rotation Matrix: Same as orthogonal transforms, but

$$- \begin{bmatrix} x_2 \\ y_2 \end{bmatrix} = \begin{bmatrix} r_{11} & r_{12} \\ r_{21} & r_{22} \end{bmatrix} \begin{bmatrix} x_1 \\ y_1 \end{bmatrix} + \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

-

$$r_{11}r_{12} + r_{21}r_{22} = 0$$

$$r_{11}r_{11} + r_{21}r_{21} = 1$$

$$r_{12}r_{12} + r_{22}r_{22} = 1$$

$$\det(R) = 1$$

- Can be used for rotations.
- Preserves linearity AND distance.

- Orientation: altitude and (heading or yaw)

- Pose: position and orientation

$$- \text{2D: } \begin{bmatrix} x & y & \psi \end{bmatrix}^T$$

$$- \text{3D: } \begin{bmatrix} x & y & z & \theta & \phi & \psi \end{bmatrix}^T$$

- : Posture: pose plus some configuration

3.1.3 Homogeneous Transforms

- Pure Direction

- Points in 3D can be rotated, reflected, scaled and sheared with 3x3 matrices but not translated.
- Trick: Move to 4D
-

$$p_2 = p_1 + p_k = \begin{bmatrix} x_1 \\ y_1 \\ z_1 \\ 1 \end{bmatrix} + \begin{bmatrix} x_k \\ y_k \\ z_k \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & x_k \\ 0 & 1 & 0 & y_k \\ 0 & 0 & 1 & z_k \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ y_1 \\ z_1 \\ 1 \end{bmatrix} = trans(p_k)p_1$$

- Operators

$$- trans(u, v, w) = \begin{bmatrix} 1 & 0 & 0 & u \\ 0 & 1 & 0 & v \\ 0 & 0 & 1 & w \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

$$- rot_x(\theta) = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & \cos(\theta) & -\sin(\theta) & 0 \\ 0 & \sin(\theta) & \cos(\theta) & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

– Operating on a point v.s. operating on a direction

- The columns of the identity HT can be considered to represent the coordinate frame itself.
- Homogeneous Transforms are both operators and frames. They can be both the things that operate on other things and things operated upon.

3.1.4 Euler Angles

- Can define rotation relative to axes of a frame
- Euler angles have trouble rotating about two or more axes
- Many euler angles map to one rotation (gimbal lock)

3.1.5 Quaternions

Chapter 4

Convexity and Optimization

4.1 Convex Optimization

- Mature field with deep mathematical foundations
- Scales well with dimensionality, solves problems with 1000s of variables
- Convex optimizers are usually really fast
- Functions are defined as $f : A \rightarrow B$, "f maps elements in the set A to elements in set B"
- Derivatives: a linear approximation to a function at a certain point
 - $Df(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$ for f from $\mathbb{R} \rightarrow \mathbb{R}$
 - Suppose $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$
 - * f is differentiable at x if there exists a matrix $Df(x) \in \mathbb{R}^{m \times n}$
 - * $Df(x)$ is called the derivative (or **Jacobian**) of the function
 - * $Df(x)_{ij} = \partial \frac{f_i(x)}{\partial x_j}$ for $i = 1 \dots m, j = 1, \dots, n$

4.2 Search for Optimization

- Consider non-convex continuous problems and problems where variables are discrete
- A graph is a set of vertices V and edges E
- Graphs capture the idea of adjacency and we can use adjacent relationships to search the graph for a certain node or path between nodes
- In optimization, adjacency between nodes can be used to determine what solutions to explore next
- e.g. n-queens
- Local search algorithms: used to search graph for best solution
 - To solve optimization problems using search:

- * Define set of possible solutions (nodes)
- * Define adjacency between solutions (edges)
- * Define cost/value/fitness function for nodes
- Descent methods are a form of local search algorithms
- Hill climbing: consider next possible moves and pick one that improves the cost function the most
 - * Drawbacks: depending on initial state, can get stuck in local optima
 - * Can try running algorithm some number of times with random start state. If you run enough times, you will get the answer (in the limit). Takes a lot of time, no guarantees on when to terminate.
- Simulated Annealing
 - * Explicitly inject variability into search process
 - * More variability at beginning of search and decrease this over time (don't want to move away from good solution)
 - * Using a temperature schedule
- Evolutionary Algorithms
 - Genetic algorithms: inspired by process of evolution in nature
 - Operators:
 - * Crossover: new state generated from two parent states
 - * Mutations: Randomly change component of state

Algorithm 1: Genetic algorithms

- 1 1. Initialize population (k random states);
 - 2 2. Select a set of parents from population for mating (based on fitness);
 - 3 3. Generate children via crossover of parents;
 - 4 4. Mutation (add randomness to children);
 - 5 5. Evaluate fitness of children;
 - 6 6. Replace worst parent with children;
 - 7 7. Repeat from step 2
-

4.3 Search for a Path (used in motion planning)

- Formulating a path search problem
 - State space
 - Successor Function:
 - Actions
 - Action Cost
 - Goal Test

- Tree Search Algorithms
 - Completeness: does it always find a solution if one exists?
 - Optimality: does it always find the least-cost solution?
 - two types of complexity:
 - * Time complexity
 - * Space complexity
 - Measured in terms of
 - * b : maximum branching factor of search tree
 - * d : depth of least-cost solution
 - * m : maximum depth of state space
 - Breadth-first search
 - Depth-first search
 - Best-first search
 - A* Search
 - Variants of A*
 - * Dynamic A* (D*), Lifelong Planning A*, Anytime Repairing A*

Chapter 5

Motion Planning

- Motion Planning is the automatic generation of motion / path for a robot that does not collide with obstacles.
- Path planning is global search for path to goal whereas obstacle avoidance ("local navigation") is a reactive method
 - Exact algorithms
 - * Either find a solution or prove one doesn't exist
 - * Computationally expensive
 - * Unsuitable for high-dimensional spaces
 - Discrete search
 - * Divide space into grid, use A* search
 - * Unsuitable for high-dimensional spaces
 - Sampling-based planning
 - * Sample the C-space, construct path from samples
 - * Good for high-dimensional spaces
 - * Weak completeness and optimality guarantees

5.1 Methods

- Visibility graph
 - Continuous representation (configuration space formulation)
 - Discretization (random sampling)
 - Graph searching (BFS, DFS, A*)
- A visibility graph is a graph such that
 - nodes: q_{init} , q_{goal} or obstacle vertex

- edges: edge exists between nodes u and v if the line segment between u and v is an obstacle edge or does not intersect the obstacles
- Algorithm
- Computational Efficiency
- Cell decomposition: decompose free space into simple cells and represent connectivity of free space by adjacency graph of these cells
-
- Potential field: define potential function over free space that has global minimum at goal and follow the steepest descent of the potential function

5.2 Configuration Space

- Open set - a set with no boundary. Every point in the set has an open neighborhood which is also in the set.
- Closed set - a set with a boundary. A closed set is a complement of some open set and vice versa.
- A set X is called a topological space if there is a collection of open subsets of X such as
 - the union of any number of open sets is an open set
 - the intersection of a finite number of open sets is an open sets
 - both X and \emptyset are open sets
- Two topological spaces X and Y are **homeomorphic** if there is a bijective function $f : X \rightarrow Y$ and both f and f^{-1} are continuous.
- Homeomorphisms can not add or remove holes.
- Common Topological Spaces: the real numbers (\mathbb{R}^1), the unit circle (\mathbb{S}^1)
- Can make more complex spaces using the Cartesian product.
- $\mathbb{R}^1 \times \mathbb{S}^1 =$ a hollow cylinder
- The **configuration** of a moving object is a specification of the position of every point on the object
 - A configuration q is usually expressed as a vector of the DOF of the robot
- The **configuration space** C is the set of all possible configurations. Usually this is a topological space. A configuration q is a point in C .
- The **dimension of a configuration space** is the minimum number of DOF needed to specify the configuration of the object completely.
- An **articulated** object is a set of rigid bodies connected by joints.
- A **path** in C is a continuous curve connecting two configurations q_{start} and q_{goal} .
- A **trajectory** is a path parameterized by time.

- A configuration q is collision-free if the robot placed at q does not intersect any obstacles in the workspace.
- The free space C_{free} is a subset of C containing all free configurations.
- A configuration space obstacle C_{obs} is a subset of C that contains all configurations where the robot collides with workspace obstacles or with itself.
- Minkowski sum [insert image]

5.3 Sampling-based Planning

- How do we plan in **high-dimensional** C-spaces?
- Exact methods either find a solution or prove none exists
- They require computing C-space obstacles which are very computationally expensive!
- Discrete search run-time and memory requirements are sensitive to branching factors (# of successors)
- Number of sessions depends on dimension, n-dimensional 8-connected space has $3^n - 1$ successors
- In sampling-based planning, instead of systematically-discretizing the C-space, take samples in the C-space and use them to construct a graph.
- Advantages
 - Don't need to discretize
 - Don't need to explicitly represent C-space
 - Easy to sample high-dimensional spaces
- Disadvantages
 - Probability of sampling an area depends on the area's size, hard to sample narrow passages
 - No guarantees on completeness / optimality
- Probabilistic Roadmap (PRM) - multi-query algorithms because roadmap can be reused if environment and robot haven't changed between queries
 - Build a roadmap of the space from sampled points and search roadmap to find a path
- 1. "Learning" Phase
 - (a) Construction step
 - i. Build roadmap by sampling random free configurations and connect them using a fast local planner
 - ii. Store configurations as nodes in a graph
 - iii. Edges of graph are paths between nodes found by local planner

- Need a distance metric to define "nearest": $D(q_1, q_2)$, use Euclidean distance
- Naive NN can be slow with 1000s of nodes, so use **kd-tree** to store nodes and do NN queries
- Kd-tree is a data structure that recursively divides the space into bins that contains points (like Oct-tree) and nearest neighbor searches through bins to find nearest point
- Local planner can be anything, but must be fast because it is called many times by the algorithm
- Easiest way is just to connect points using a straight line and check whether the line is collision free

(b) Expansion step

- You can have disconnected components that should be connected
- Expansion uses heuristics to sample more nodes to connect disconnected components
- No "right" way, this step is environment dependent

2. Query Phase

- Given start q_s and goal q_g . Connect them to the roadmap using a local planner.
- Then search the graph G to find the shortest path between q_s and q_g using A*, Dijkstra's, etc.

Algorithm 2: Path shortening / smoothing

```

1 for  $i = 0$  to  $maxiterations$  do
2   | pick two points  $q_1$  and  $q_2$  on the path randomly;
3   | try to connect them with a line segment;
4   | if successful, replace path between  $q_1$  and  $q_2$  with the line segment;
5 end

```

3. PRM Failure Modes

- Cannot connect q_s and q_g to any nodes in the graph
- Cannot find a path in the graph but path is possible

4. PRM issues

- Uniform random sampling misses narrow passages
- Exploring whole space, but all we want is a path

5. Sampling strategies

- Gaussian sample
- Bridge sample

- * Sample a q_1 in collision
- * Sample a q_2 in neighborhood of q_1 with some prob distribution
- * If q_2 is in collision, get the midpoint of (q_1, q_2)
- * Check if midpoint is in collision and if not add it as a node
- Rapidly-exploring Random Trees (RRTs) : single-query method
 - Build a **tree** instead of a graph**
 - The tree grows in C_{free}
 - Like PRM captures some connectivity, but unlike PRM it only explores what is connected to q_{start}
 - RRT Goal Biasing
 - * Bias RRTs towards goal to produce a path
 - * When generating a random sample, with some probability pick the goal instead of random node
 - RRT Extension Types
 - * RRT-Extend: Take one step towards a random direction
 - * RRT-Connect: Step towards random sample until it is either reached or you hit an obstacle
 - * BiDirectional RRTs: grow tree from both start and goal
 - * RRT produces bad paths, must perform path smoothing (ALWAYS)

Algorithm 3: Naive Tree algorithm

```

1  $q_{node} = q_{start};$ 
2 for  $i = 1$  to  $num\_samples$  do
3    $q_{rand} = \text{sample near } q_{node};$ 
4   Add edge  $e = (q_{rand}, q)$  if collision-free;
5    $q_{node} = \text{pick random node of tree};$ 
6 end
```

Algorithm 4: Build RRT

```

1 T.init( $q_{init}$ );
2 for  $k = 1$  to  $K$  do
3    $q_{rand} = \text{random\_config}();$ 
4   extend(T,  $q_{rand}$ );
5 end
```

5.4 Nonholonomic Planning

- **Holonomic constraints** depend only on configuration
- $F(q, t) = 0$

- These have to be bilateral constraints (no inequalities)
- Example: kinematics of a unicycle
 - Can move forward and backward
 - Can rotate about the wheel center
 - Can't move sideways
- **Non-holonomic** constraints are non-integrable. Thus they must contain derivatives of configuration. Sometimes called differential constraints.
- State space (configuration + velocity) vs control space (speed or acceleration, steering)
- Simple Car
 - Non-holonomic constraint: $-\dot{x}\sin(\theta) + \dot{y}\cos(\theta) = 0$
 - Essentially means you can't move sideways
- Two-point Boundary Value Problem (BVP): find a control sequence to take system from state x_i and x_g while obeying kinematic constraints
- Discrete Planning Option 1: Sequencing primitives
 - Discretize control space into primitives (pick steering angles, accelerations, velocities)
 - Disadvantage: losing full continuous completeness and discontinuous curvature
 - Choice of primitives affects completeness, optimality, and speed
- Discrete Planning Option 2: State Lattice
 - Pre-compute state lattice
 - Two methods to get lattice:
 - * Forward: using motion primitives
 - * Inverse: Use BVP solvers to find trajectories between states
 - Impose continuity constraints at graph vertices
 - Search state lattice like any graph
- Sampling-based Planning
 - Building state lattice is impractical in high dimensions
 - We are now sampling state space, not C-space!
 - Challenges
 - * Dimension of space is doubled
 - * Moving between points is harder
 - * Distance metric is unclear (worst problem)
 - RRT Non-holonomic Planning

*

Chapter 6

Kinematics

Chapter 7

Grasping

- Grasping studies how to stably make contact with objects and move them
- Definitions

7.1 Definitions

- A point contact is sometimes called a finger
- A **wrench** is a combination of force and torque applied to an object
- **Wrench space** is the space of wrenches applied to an object
 - 2D object: 3 dimensional wrench space (2 force, 1 torque)
 - 3D object: 6 dimensional wrench space (3 force, 3 torque)
- A grasp **immobilizes** an object if it can counter any wrench applied to the object. This guarantees the stability of the grasp.
- A **friction cone** is the set of forces that can be applied at a contact force without sliding on the object. Assume Coulomb friction.
 - Depends on the coefficient of friction between hand and object (μ)
 - Bigger μ implies a wider friction cone.

7.2 Form Closure

- A form closure grasp is when the object cannot move **regardless of surface friction**
- You need at least $N+1$ contacts to achieve first-order form closure, where N is the number of DOF of the object

7.3 Force Closure

- Frictional properties of the object can be used to immobilize it

- If a grasp achieves form closure, it also achieves force closure
- Intuition, need a contact force to cancel out external disturbance force. Convex hull must contain origin for there to exist such a contact force.
- For a 3D object, you only need 3 contacts to achieve force closure (as opposed to 7 for form closure)
- Force Closure Metrics
 - Popular metric: radius of largest hyper-sphere you can fit in convex hull
 - Task specific metric: using an ellipsoid instead of a hyper-sphere

Algorithm 5: Testing for force closure

- 1 Input: Contact locations;
 - 2 Output: Is the grasp in force-closure?;
 - 3 1. Approximate friction cone at each contact with a set of wrenches.;
 - 4 2. Combine wrenches from all cones to a set of points S in wrench space.;
 - 5 3. Compute the convex hull of S (smallest convex set that contains all points);
 - 6 4. If the origin is inside the convex hull, return YES. Else return NO.;
-

7.4 Searching for Force Closure Grasps

- Peter Allen et al. 2000s
 - Sample pose of hand relative to object with fingers in a pre-determined shape
 - Approach object until contact and close fingers
 - Get contact points between hand and object
 - Test these contact points for force closure
- Pre-compute grasp sets: searching for grasps is slow!
- Columbia Grasp Database

7.5 Integrating Grasping and Motion Planning

- Pre-compute grasp set offline, get force-closure scores
- Online: compute 2 scores for each grasp: Environment Clearance Score and Reachability Score
- Test grasps in order of ranking
- Recent work in grasping uses deep learning methods
- General idea
 - Generate many grasp candidates
 - Learn a quality metric that uses the point cloud data directly
 - Output highest quality grasp

Chapter 8

SVD and PCA

How do we transform the data to get rid of "unimportant" dimensions/rotations?

8.1 Definitions

- **Variance** is the measure of deviation from the mean for points in one dimension
- **Covariance** is the measure of how much each dimensions vary from the mean with respect to each other
- Covariance is measured between pairs of dimensions to see their correlation
- Magnitude of covariance is not as important as sign
- + covariance means both dimensions increase or decrease together
- - covariance means one increases while the other decreases
- covariance = 0 means the dimensions are independent of one another
- Estimate covariance matrix:
 1. Subtract the mean of the datapoints from every column of X
 2. $Q = \frac{XX^T}{n-1}$
- Eigenvalue problem
 - A: $n \times n$ matrix
 - v: $n \times 1$ non-zero vector
 - λ : scalar
 - $Av = \lambda v$
- A value of λ for which this equation has a solution is called an **eigenvalue** of A
- A v corresponding to this value of λ is called an **eigenvector** of A
- All eigenvectors of a matrix are *orthogonal* to each other
- Eigenvectors of a covariance matrix

- Eigenvectors of Q with the largest eigenvalues correspond to dimensions that have the strongest correlation in the dataset
- Eigenvectors of Q are **principle components**

8.2 Singular Value Decomposition (SVD)

- SVD decomposes any matrix M into $M = U\Sigma V^T$
- The columns of U are the eigenvectors of MM^T
- Σ is a diagonal matrix where the elements of the diagonal are the $\sqrt{\text{eigenvalues}}$ of $M^T M$ and MM^T in decreasing order of magnitude
- Columns of V are the eigenvectors of $M^T M$
- Columns of U and V are **orthonormal** meaning each column vector has unit magnitude and orthogonal to all other column vectors.

8.3 Principle Component Analysis (PCA)

- We care about the variance of the data
- High variance = high importance
- PCA is a technique used
 - Remove rotation in a dataset
 - Reduce dimensionality of a dataset
- PCA computes linear transformation that chooses a new coordinate system for the data set such that the greatest variance by any projection of the data set comes to lie on the first axis (called the **first principal component**)

Algorithm 6: PCA

- 1 1. Given dataset X ;
 - 2 2. Compute mean of X ;
 - 3 3. $X = X - \mu$ (subtract mean from every point in X);
 - 4 4. Compute covariance Q of X ;
 - 5 5. Take the SVD of $Q = U\Sigma V^T$;
 - 6 6. $X_{new} = V^T X$;
-

8.3.1 Limitations of PCA

- PCA is sensitive to the scaling of the variables
- ** need to review this part

8.3.2 Applications of PCA

- Eigenfaces: analysis of database of face images

- Instead of using all the pixel values, we can represent a face as a weighted combination of eigenfaces
- SVD takes a long time to run for high-dimensional data
- Can be used for video compression
- Many high-dimensional datasets have hidden low-dimensional structure

Chapter 9

Probabilistic Models

9.1 Probability Basics

9.1.1 Discrete Random Variables

- X denotes a random variable and it can take on a countable number of values in $\{x_1, x_2, \dots, x_n\}$
- $P(X = x_i)$ is the probability that the random variable X takes on value x_i
- $P(\dots)$ is called the **probability mass function**
- e.g. $P(\text{Room}) = \langle 0.7, 0.2, \dots, 0.02 \rangle$

9.1.2 Continuous Random Variables

- X takes on a value in the continuum
- $P(X = x)$ is the *probability density function*
- $P(x \in (a, b)) = \int_a^b P(x)dx$

9.1.3 Axioms of Probability Theory

- $0 \leq P(a) \leq 1$
- $P(\text{true}) = 1$ and $P(\text{false}) = 0$
- $P(a \vee b) = P(a) + P(b) - P(a \wedge b)$

9.1.4 Joint and Conditional Probability

- $P(X = x \wedge Y = y) = P(x, y)$
- If X and Y are **independent** then $P(x, y) = P(x)P(y)$
- $P(x|y) = \frac{P(x,y)}{P(y)}$
- $P(x, y) = \frac{P(x|y)}{P(y)}$

- If X and Y are **independent** then $P(x|y) = P(x)$
- $P(x, y|z) = P(x|z)P(y|z)$ means that x and y are **conditionally independent**
- If I know z , I don't need to know x to compute the probability of y .

9.1.5 Law of Total Probability (Discrete)

- $\sum_x P(x) = 1$
- $P(x) = \sum_y P(x, y) = \sum_y P(x|y)P(y)$

9.2 Bayes Rule

- $P(x|y) = \frac{P(y|x)P(x)}{P(y)} = \frac{\text{likelihood} \cdot \text{prior}}{\text{evidence}}$
- Usually $P(y)$ is difficult to compute, so use normalization trick
- $P(x|y) = \eta P(y|x)P(x)$ where $\eta = \frac{1}{\sum_{x \in X} P(y|x)P(x)}$

9.2.1 Casual and Diagnostic Reasoning

- Suppose a robot wants to determine probability of a door being open
- It obtains measurement z . What is the $P(open|z)$
- $P(open|z)$ is **diagnostic** and $P(z|open)$ is **causal**

9.2.2 Conditional Independence Example

- Consider three variables: *RobotLocation*, *GPSEstimate*, *LandmarkEstimate*
- *GPSEstimate* and *LandmarkEstimate* are NOT independent, $P(GPSEstimate|LandmarkEstimate) \neq P(GPSEstimate)$
- *GPSEstimate* and *LandmarkEstimate* are conditionally independent given *RobotLocation*
- If I know the robot's location, then I can compute the landmark estimate without knowing the GPS estimate

9.3 Bayes Net

- Encode conditional independence relationships in a Bayes Net. Used to describe cause-effect relationships
- Directed and acyclic graph
 - Nodes represent random variables
 - Edges represent conditional dependencies
 - Nodes that are not connected are conditionally independent of each other
 - Node is associated with a probability function $P(X_i|Parents(X_i))$, this is defined by a conditional probability table (CPT)

- Inference

9.3.1 Markov Random Fields

- Graph is undirected and may be cyclic

9.3.2 Conditional Random Fields

- Undirected graphical model whose nodes can be divided into exactly two disjoint sets:
 - X: the input variables
 - Y: the observed and output variables
- Used to model the conditional distribution: $P(Y|X)$
- CRFs can be used for object recognition and image segmentation

9.4 Learning a probabilistic model

-

Chapter 10

Filters

10.1 Bayes Filter

- Bayes Filter accounts for both robot state and perception data
- To use Bayes filter, the state must be discrete, usually represented by a grid
- Each grid cell, contains the **belief** (probability that the true state of the system is x_t)

Algorithm 7: Discrete Bayes Filter Algorithm

```
1 Inputs: Bel(x), d;  
2  $\eta = 0$ ;  
3 if  $d$  is a perceptual data item  $z$  then  
4   for all  $x$  do  
5      $Bel'(x) = P(z|x)Bel(x)$ ;  
6      $\eta = \eta + Bel'(x)$ ;  
7   end  
8   for all  $x$  do  
9      $Bel'(x) = \eta^{-1}Bel'(x)$ ;  
10  end  
11 end  
12 else if  $d$  is an action data item  $u$  then  
13   for all  $x$  do  
14      $Bel'(x) = \sum_{x'} P(x|u, x')Bel(x')$ ;  
15   end  
16 end  
17 Return  $Bel'(x)$ 
```

10.2 Kalman Filter

- Kalman filter used when state space is continuous variables
- They key idea is to represent everything with *gaussians*
- Univariate and multivariate gaussians

- We stay in "Gaussian world" as long as we start with Gaussians and perform only linear transformations
- Estimate state x of a *discrete-time* controlled process governed by linear stochastic difference equation:

$$x_t = A_t x_{t-1} + B_t u_t + \epsilon_t$$

and sensor measurement

$$z_t = C_t x_t + \delta_t$$

where:

x_t = current state

A_t = matrix describing how state changes from $t - 1$ to t without controls

B_t = matrix that describes how control u_t changes the state from $t - 1$ to t

C_t = matrix that describes how to map state x_t to an observation z_t

ϵ_t = process noise normally distributed with covariance R_t

δ_t = measurement noise normally distributed with covariance Q_t

Algorithm 8: Kalman Filter

- 1 Prediction: use dynamics to predict what will happen;
 - 2 $\bar{\mu}_t = A_t \mu_{t-1} + B_t u_t$;
 - 3 $\bar{\Sigma}_t = A_t \Sigma_{t-1} A_t^T + R_t$;
 - 4 Correction: use sensor measurement to correct prediction;
 - 5 $K_t = \bar{\Sigma}_t C_t^T (C_t \bar{\Sigma}_t C_t^T + Q_t)^{-1}$;
 - 6 $\mu_t = \bar{\mu}_t + K_t (z_t - C_t \bar{\mu}_t)$;
 - 7 $\Sigma_t = (I - K_t C_t) \bar{\Sigma}_t$;
 - 8 **return** μ_t, Σ_t
-

- Comments:
 - Highly efficient: only need to compute matrix multiplication
 - Optimal for linear Gaussian systems, but most robotics systems are nonlinear

10.3 Extended Kalman Filter (EKF)

- Most robotics problem deal wth nonlinear dynamics and sensors

$$x_t = g(u_t, x_{t-1})$$

$$z_t = h(x_t)$$

- EKF trick: use a *local linear approximation* by computing the Jacobians of g and h

$$x_t = g(u_t, x_{t-1}) \approx g(u_t, \mu_{t-1}) + G_t (x_{t-1} - \mu_{t-1})$$

$$z_t = h(x_t) \approx h(\bar{\mu}_t) + H_t (x_t - \bar{\mu}_t)$$

where

$$G_t = \frac{\partial g(u_t, \mu_{t-1})}{\partial x_{t-1}}$$

$$H_t = \frac{\partial h(\bar{\mu}_t)}{\partial x_t}$$

- Comments
 - Highly efficient
 - Not optimal, because it is an approximation of the nonlinear function
 - Can diverge if nonlinearities are large (e.g. close to beacon)
 - Everything must be a Gaussian
 - Cannot be used if transition is non-linear, e.g. multimodal distributions

10.4 Unscented Kalman Filter

- Key idea:
 1. Sample a set of sigma points from Gaussian distribution
 2. Pass sigma points through function
 3. Re-estimate Gaussian

Algorithm 9: Unscented Kalman Filter

- 1 \mathcal{X}_{t-1} = Compute sigma points using μ_{t-1} and Σ_{t-1} ;
 - 2 $\mathcal{X}_t^* = g(u_t, \mathcal{X}_{t-1})$ // Apply dynamics ;
 - 3 $\bar{\mu}_t = \sum_{i=0}^{2n} w_m^{[i]} \bar{\mathcal{X}}_t^{*[i]}$ // Estimate new mean from sigma points;
 - 4 $\bar{\Sigma}_t = \sum_{i=0}^{2n} w_c^{[i]} (\bar{\mathcal{X}}_t^{*[i]} - \bar{\mu}_t)(\bar{\mathcal{X}}_t^{*[i]} - \bar{\mu}_t)^T + R_t$;
 - 5 $\bar{\mathcal{X}}_t$ = Compute sigma points using $\bar{\mu}_t$ and $\bar{\Sigma}_t$;
 - 6 $\bar{\mathcal{Z}}_t = h(\bar{\mathcal{X}}_t)$ // Apply sensor model;
 - 7 TODO;
-

- Comments
 - Highly efficient
 - Better approximation than EKF: accurate in first two terms of Taylor expansion
 - Derivative-free, does not require any Jacobians

10.5 Particle Filter

- No need to make assumptions of distributions unlike Kalman Filter which assumes that all error is Gaussian
- Sample from the implicit distribution of the state at any given time
- At time t , distribution is represented by the M samples of the robot's states

$$X_t = x_t^{[1]}, x_t^{[2]}, \dots, x_t^{[M]}$$

$$W_t = w_t^{[1]}, w_t^{[2]}, \dots, w_t^{[M]}$$

- If starting location is know, can initialize particles around start, else scatter particles through the environment

Algorithm 10: Particle Filter

```

1  $X_0 = \text{Sample } M \text{ particles from } P(X_0);$ 
2  $t = 0;$ 
3 while True do
4    $t++;$ 
5    $u_t = \text{action}();$ 
6    $z_t = \text{sensor}();$ 
7    $S_t = X_t = \{\};$ 
8   for  $m = 1$  to  $M$  do
9     sample  $x_t^{[m]} \sim p(x_t|u_t, x_{t-1}^{[m]} \text{ from } X_{t-1});$ 
10     $w_t^{[m]} = p(z_t|x_t^{[m]});$ 
11     $S_t = S_t \cup (x_t^{[m]}, w_t^{[m]});$ 
12  end
13  for  $m = 1$  to  $M$  do
14    draw  $i$  with probability  $\sim w_t^{[i]};$ 
15    add  $x_t^{[i]}$  from  $S_t$  to  $X_T;$ 
16  end
17 end

```

- Sampling strategies
 - Need to preserve diversity of particles, inject random particles
 - Low-variance sampling
 - Stratified sampling
- Comments
 - Represent distribution with a set of particles
 - Approximate arbitrary probability distributions
 - More particles = better approximation but more computation cost

Chapter 11

MDP and POMDPs