

# Anthony Liang

[aliang8.github.io](https://aliang8.github.io) ◇ [anthony.liang@usc.edu](mailto:anthony.liang@usc.edu) ◇ [linkedin.com/in/ant-liang](https://www.linkedin.com/in/ant-liang)

## RESEARCH INTERESTS

My research focuses on scalable robot learning, with an emphasis on data-efficient human feedback (gaze, demonstrations, preferences, interventions) and reward foundation models that provide generalizable reward functions across tasks, embodiments, and environments. I also work on core components for embodied intelligence, including vision-language-action models for low-level control, trace models for structured planning, and world models for scene understanding and predictive reasoning. In parallel, I am investigating RL post-training for LLM-based web agents to improve reliability in sequential UI tasks and long-horizon interactions. This line of work parallels robot learning through shared challenges in reinforcement learning and human feedback.

## EDUCATION

<b>University of Southern California</b> , Los Angeles, CA <i>Ph.D.</i> in Computer Science (Co-advised by Erdem Biyik and Stephen Tu)	Aug 2021 - Present GPA: 3.9/4.0
<b>University of Michigan Rackham Graduate School</b> , Ann Arbor, MI <i>M.S.</i> in Robotics (Advisor: Honglak Lee)	Aug 2017 - 2021 GPA: 4.0/4.0
<b>University of Michigan</b> , Ann Arbor, MI <i>B.S.E.</i> (Advisor: Honglak Lee)	Aug 2017 - 2021 GPA: 3.65/4.0

## RESEARCH EXPERIENCE

<b>LiRA Lab and Statistical Learning Lab</b> , USC <i>Ph.D. Student, PI: Erdem Biyik, Stephen Tu</i>	Aug 2021 - Present
- Robot learning with multimodal human feedback, sample-efficient RL, learning from unlabelled videos	
<b>Intelligent Robot Lab</b> , Carnegie Mellon <i>Visiting Researcher</i> with Changliu Liu	May 2020 - May 2021
- Hierarchical RL for safe control of autonomous vehicles in dynamic environments	
<b>Deep Learning Lab</b> , University of Michigan <i>Research Intern</i> with Honglak Lee	Jan 2019 - May 2021
- Sample-efficient RL for embodied task learning	

## PROFESSIONAL EXPERIENCE

<b>Google DeepMind</b> <i>Research Intern</i> with Kalpesh Krishna, Jacob Eisenstein	May 2025 - Sept 2025 NYC/Seattle
- RL post-training for reasoning models	
<b>Google Research</b> <i>Research Intern</i> with Chih-Wei Hsu, Yinlam Chow, Guy Tennenholtz, Craig Boutilier	May 2023 - Feb 2024 Remote NYC
- Bayesian RL for Markov Decision Processes with gradually changing latent dynamics	
<b>Meta AI - Multimodal Learning Team</b> <i>Research Intern</i> with Paul Crook and Andrea Madatto	May 2022 - Aug 2022 Redmond, WA
- Fine-tuning large language models for task-oriented dialogue generation	
<b>Amazon Science</b> <i>Applied Science Intern</i> with Thiago Mosquiero	May 2021 - Aug 2021 Seattle, WA
- Collaborative filtering for recommending new brands and products to consumers	
<b>Invisible.ai</b> <i>AI Research Intern</i>	May 2020 - Aug 2020 Remote
- Improving computer vision models for real-time object detection and tracking for industrial processes	

**Google Ads**  
*Software Engineering Intern*

May 2019 - Aug 2019  
*Mountain View, CA*

**Luminar Technologies**  
*Software Engineering Intern*

May 2018 - Aug 2018  
*Palo Alto, CA*

## PUBLICATIONS

---

- [C8] [Anthony Liang](#), Jonathan Berant, Adam Fisch, Abhimanyu Goyal, Kalpesh Krishna, Jacob Eisenstein. "Plantain: Plan-Answer Interleaved Reasoning" *Preprint*
- [C7] Mihir Rao\*, [Anthony Liang](#)\*, Abrar Anwar\*, Rudy Corona, Giscard Biamby, Dantong Liu, Roei Herzig, Trevor Darrell. "In-Context Learning Trace Model for Robot Learning" *In preparation*
- [C6] Matthew Hong\*, [Anthony Liang](#)\*, Kevin Kim, Harshitha Rajaprakash, Jesse Thomason†, Erdem Biyik†, Jesse Zhang†. "HAND Me the Data: Fast Robot Adaptation via Hand Path Retrieval" *In submission*
- [C5] Jaiv Doshi\*, [Anthony Liang](#)\*, Yigit Korkmaz, Erdem Biyik. "RHODES: Reducing Human Oversight via Disagreement and Exploration for Safe Reinforcement Learning" *Preprint*
- [C4] [Anthony Liang](#)\*, Pavel Czempin\*, Matthew Hong, Yutai Zhou, Erdem Biyik, Stephen Tu. "CLAM: Continuous Latent Action Models for Robot Learning from Unlabeled Demonstrations" *In submission*
- [C3] [Anthony Liang](#), Chih-wei Hsu, Yinlam Chow, Guy Tennenholz, Erdem Biyik, Craig Boutilier. "DynaMITE-RL: A Dynamics Model for Improved Temporal Meta Reinforcement Learning", *International Conference on Machine Learning (ICML) AutoRL Workshop 2024, Conference on Neural Information Processing Systems (NeurIPS) 2024*
- [C2] [Anthony Liang](#), Jesse Thomason, Erdem Biyik. "ViSaRL: Visual Reinforcement Learning Guided By Human Saliency", *Spotlight talk at ICRA Pretraining for Robotics (PT4R) Workshop, IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS) 2024*
- [C1] Wilka Carvalho, [Anthony Liang](#), Kimin Lee, Sungryull Sohn, Honglak Lee, Richard L. Lewis, Satinder Singh. "Reinforcement Learning for Sparse-Reward Object-Interaction Tasks in a First-person Simulated 3D Environment", *International Joint Conferences on Artificial Intelligence (IJCAI) 2021*
- [T5] [Anthony Liang](#), Yigit Korkmaz, Jiahui Zhang, Jesse Zhang, Abrar Anwar, Sidhant Kaushik, Yufei Wang, Yu Xiang, David Held, Dieter Fox, Abhishek Gupta, Stephen Tu, Erdem Biyik. "SPUR: Scaling Reward Learning from Human Demonstrations". *CoRL Eval&Deploy Workshop 2025*
- [T4] [Anthony Liang](#), Pavel Czempin, Yutai Zhou, Stephen Tu, Erdem Biyik. "In-Context Generalization to New Tasks From Unlabeled Observation Data", *ICML In-Context Learning Workshop 2024*
- [T3] Ishika Singh, [Anthony Liang](#), Mohit Shridhar, Jesse Thomason. "Self-Supervised 3D Representation Learning for Robotics", *ICRA Pretraining for Robotics (PT4R) 2023*
- [T2] [Anthony Liang](#), Ishika Singh, Karl Pertsch, Jesse Thomason. "Transformer Adapters for Robot Learning", *CoRL Workshop on Pretraining for Robot Learning 2022*
- [T1] Wilka Carvalho, [Anthony Liang](#), Kimin Lee, Sungryull Sohn, Richard L. Lewis, Satinder Singh, Honglak Lee. "ROMA: A Relational, Object-Model Learning Agent for Sample-Efficient Reinforcement Learning", *ICML Workshop on Object-Oriented Learning 2020*

## TEACHING

---

**University of Southern California**  
CSCI 699: Robot Learning  
CSCI 499: Natural Language for Interactive AI

Fall 2024  
Fall 2022

**University of Michigan, Ann Arbor**  
EECS 442: Computer Vision  
EECS 498: Algorithmic Robotics  
EECS 504: Graduate Computer Vision  
EECS 280: Introduction to Programming and Data Structures  
Summer STEM Institute Research Mentor

Winter 2021  
Fall 2020  
Winter 2020  
Fall 2018 - Fall 2019  
Summer 2021

## HONORS AND AWARDS

---

- NSF Graduate Research Fellowship Honorable Mention

2020

## SERVICES

---

- ICML 2023, 2024, 2025
- ICLR 2024, 2025
- ICRA 2024, 2025, 2026
- NeurIPS 2022, 2023, 2024, 2025
- HRI 2025
- RA-L
- AAAI 2025
- RO-MAN 2024

## STUDENT MENTORSHIP

---

- |   |   |
|---|---|
| • Mihir Rao (Undergrad, Berkeley)                       | Trace VLAs                                |
| • Ziyi Liu (Masters)                                    | Preference-based RL                       |
| • Sankalp (Sunny) Agrawal (Undergrad, USC SURE Program) | Meta-RL with task descriptors             |
| • Shreya Ramanujam (Undergrad, IIT)                     | Gaze for robot teleoperation              |
| • Matthew Hong (Masters, USC)                           | Learning from unlabelled data, RLHF       |
| • Dhanush Kumar Penmetsa (Masters, USC)                 | Gaze for robot teleoperation              |
| • Jaiv Doshi (Undergrad, USC)                           | Human-intervention reinforcement learning |
| • Yixi Quan (Undergrad, USC)                            | Pretrained Video-LLMs for Robot Learning  |
| • Junu Song (Undergrad, USC CURVE Fellowship)           | Real-world robot navigation               |